

Cloud Engineering Intern

Requirements –

- This assignment can be done in AWS or in GCP with equivalent resources (Storage, Cloud SQL, Cloud Functions, etc).
- Deployment of all the cloud resources (S3, Lambda, Database) must be automated using any Infrastructure as Code tool (CDK, Pulumi, Cloudformation, Terraform, etc) of your choice.
- API can be hosted on any platform of your choice.
- The database can be deployed in a public subnet for this assignment.
- Please don't store any public endpoint details or credentials on a public GitHub repo.
- Candidates should send a video of the working functionality
- The API code should be written in NodeJS or typescript and with proper linting.
- API endpoint should be protected by basic authentication, API key, or OAuth
- The lambda to API calls should happen in batches of 100. So, for a CSV with 1000 records, 10 API calls should be made

1. Develop a data API that gets the JSON input and stores the data in a database hosted on a cloud provider. Following is an example of the input -

```
{  
  "id": "12345"  
  "name": "John",  
  "surname": "Doe",  
  "dob": "11/11/2011",  
  "gender": "male"  
}
```

CloudFormation Template for launching AWS RDS MySQL:

AWSTemplateFormatVersion: '2010-09-09'

Description: This stack creates an RDS MySql Instance

Parameters:

MasterUsername:

Description: Database administration name.

Type: String

Default: rdsroot

MasterUserPassword:

NoEcho: 'true'

Description: Database administration password.

Type: String

MinLength: '8'

AllowedPattern: "[a-zA-Z0-9!~]*"

ConstraintDescription: Must only contain upper and lowercase letters and numbers

TcpPort:

Description: Enter RDS Listening TCP Port number.

Type: Number

Default: '3306'

DatabaseName:

Type: String

Resources:

MyRDSDBInstance:

Type: AWS::RDS::DBInstance

DeletionPolicy: Snapshot

Properties:

AllocatedStorage: "20"

DBName: !Ref DatabaseName

DBInstanceClass: db.t2.micro

DBInstanceIdentifier: !Join ["-",["MyDbInstance",!Ref "AWS::Region"]]

Engine: MySQL

MasterUsername: !Ref MasterUsername

MasterUserPassword: !Ref MasterUserPassword

Port:

Ref: TcpPort

PubliclyAccessible: True

StorageEncrypted: 'false'

StorageType: gp2

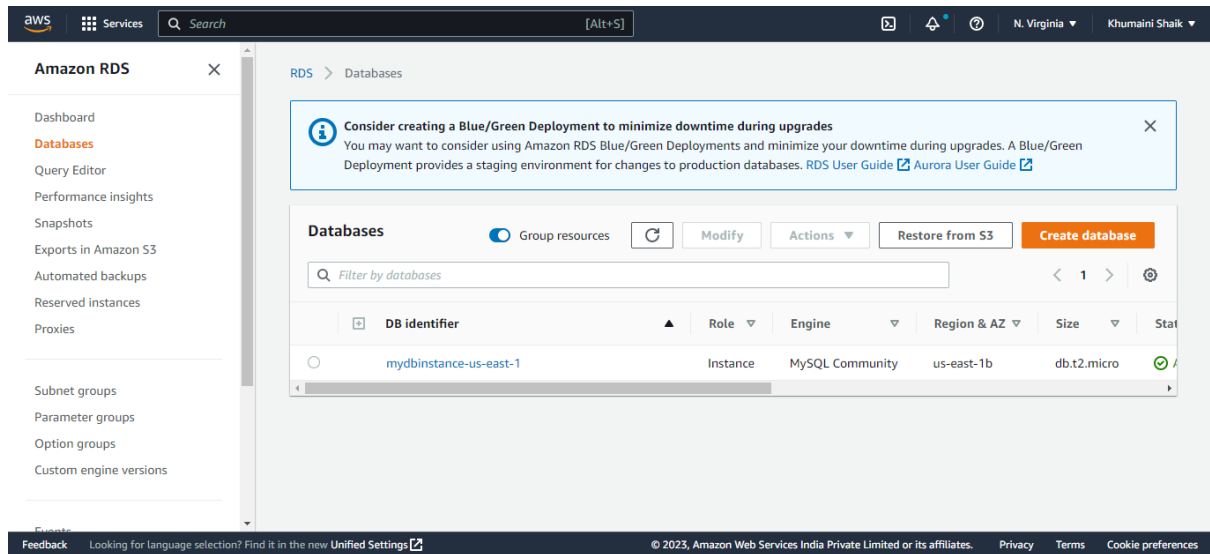
AvailabilityZone: !Select [1,!GetAZs ""]

Tags:

- Key: Owner

Value: Abdul Khumaini

Database Created:



Database Endpoint:

Connectivity & security

Endpoint & port

Endpoint

mydbinstance-us-east-
1.cc7otgq3h7hn.us-east-
1.rds.amazonaws.com

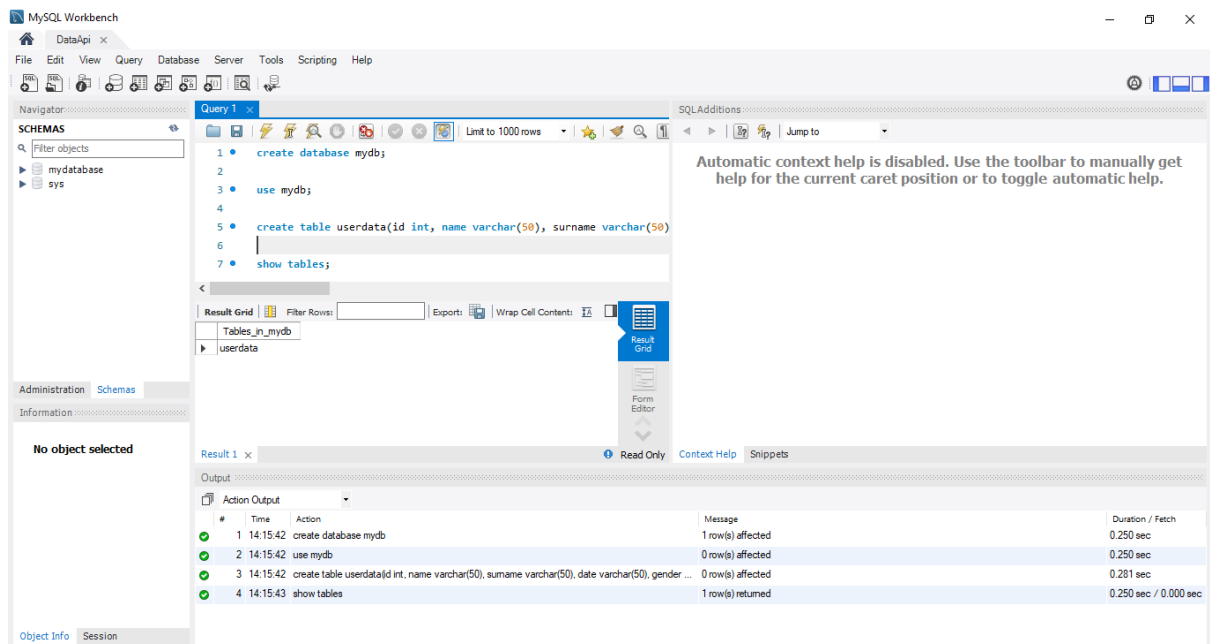
Connection created through MySQL Workbench and required sql queries are executed:

create database mydb;

use mydb;

create table userdata(id int, name varchar(50), surname varchar(50), date varchar(50), gender varchar(10));

show tables;



API to get JSON data and store it in AWS RDS, API is enabled with basic authentication to protect against attacks:

```
const express = require('express');
```

```
const app = express();
```

```
const mysql = require('mysql');
```

```
const basicAuth = require('basic-auth');
```

```
async function dbcon(req,res){
```

```
  try{
```

```
    var connection = mysql.createConnection({
```

```
      host: 'mydbinstance-us-east-1.cc7otgq3h7hn.us-east-1.rds.amazonaws.com',
```

```
      user: 'admin',
```

```

        password: 'admin1234',
        database: 'mydb',
        port:3306
    });
    connection.connect((err) => {
        if (err) {
            console.error('Database connection failed: ' + err.stack);
            return;
        }
        console.log('connected!!');
        var que = `insert into userdata
values(${req.body.id},'${req.body.name}','${req.body.surname}','${req.body.dob}','${req.bo
dy.gender}')`;
        connection.query(que,(err,res)=>{
            if(err){
                console.error(err);
            }
            console.log('1 record inserted');
        })
    });
}
catch(error){
    console.error(error);
}
}

```

```

var auth = function (req, res, next) {
    var user = basicAuth(req);
    if (!user || !user.name || !user.pass) {
        res.set('WWW-Authenticate', 'Basic realm=Unauthorized');
        res.sendStatus(401);
    }
}

```

```

    return;
  }

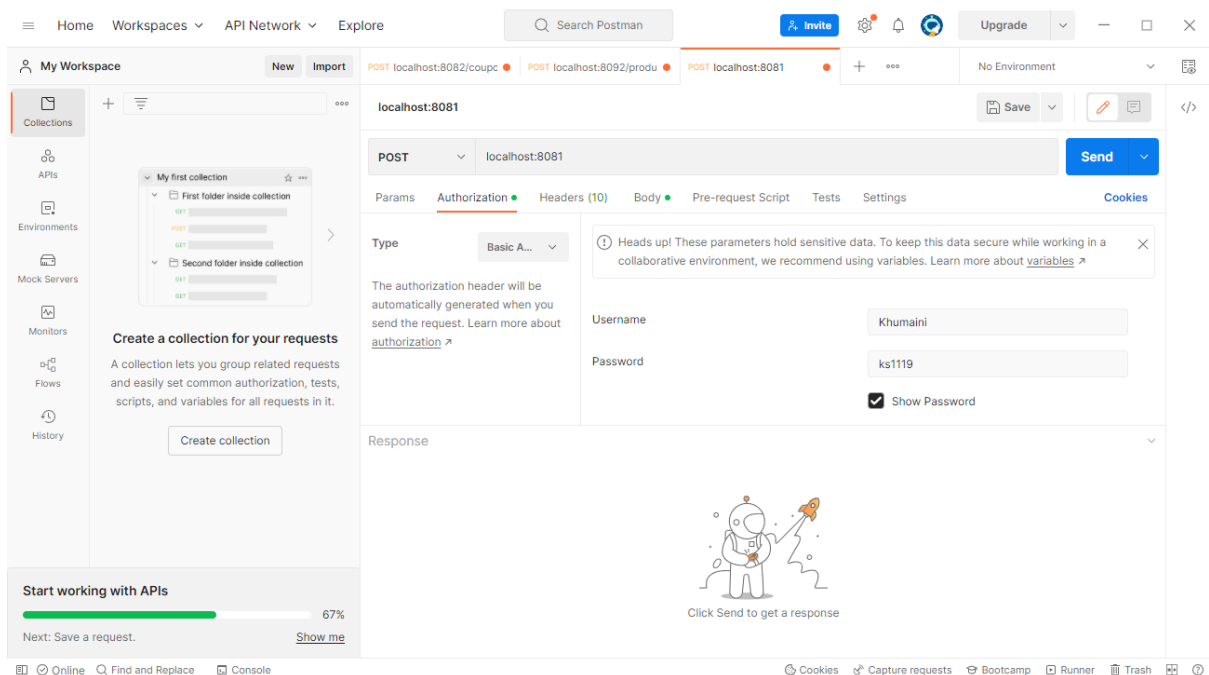
  if (user.name === 'Khumaini' && user.pass === 'ks1119') {
    next();
  } else {
    res.set('WWW-Authenticate', 'Basic realm=Unauthorized');
    res.sendStatus(401);
    return;
  }
}

app.use(express.json());
app.post('/', auth, (req, res) => {
  res.write('{ authenticated:true }\n');
  dbcon(req,res);
  res.end('1 record inserted');
});

app.listen(8081, () => console.log("Server started at 127.0.0.1://8081"));

```

Authentication:



API testing:

The screenshot shows the Postman API client interface. The top navigation bar includes 'Home', 'Workspaces', 'API Network', and 'Explore'. The main workspace is titled 'My Workspace' and shows a collection of requests. The selected request is a POST to 'localhost:8081'. The request body is a JSON object:

```
{  "id": "12345",  "name": "John",  "surname": "Doe",  "dob": "11/11/2011",  "gender": "male"}
```

. The response is a 200 OK status with a body containing

```
{  "authenticated": true,  "1 record inserted"}
```

. The interface also shows a sidebar with 'Collections', 'APIs', 'Environments', 'Mock Servers', 'Monitors', 'Flows', and 'History'. A 'Create collection' button is visible. At the bottom, there are icons for 'Online', 'Find and Replace', 'Console', 'Cookies', 'Capture requests', 'Bootcamp', 'Runner', 'Trash', and a help icon.

Server Logs:

```
E:\CloudEngineeringInternship>node restapi.js
Server started at 127.0.0.1://8081
connected!!
1 record inserted
```

Database:

The screenshot shows a database query result. The query is `select * from userdata;`. The result is displayed in a table with 6 columns: id, name, surname, date, and gender. The table contains one record with the following values: id: 12345, name: John, surname: Doe, date: 11/11/2011, gender: male. The interface also shows a 'Result Grid' button, a 'Filter Rows' input, and an 'Export' button.

id	name	surname	date	gender
12345	John	Doe	11/11/2011	male

2. Create an S3 bucket on which a user can manually upload CSV files from the AWS console. Following is the sample of the CSV -

id,name,surname,dob,gender

12357,John,Doe,11/11/2011,male

12356,Doe,Jane,01/01/2001,female

12354,Jes,Dam,11/12/2013,male

- Once the file is uploaded, it should create an event that should trigger a lambda function.
- The lambda function will read the uploaded file, parse the content of that file and post the data API.

Following is the high-level sequence diagram –



CloudFormation Template to launch S3, Lambda and to configure the EventNotification (Trigger) between them whenever a new object is created:

AWSTemplateFormatVersion: '2010-09-09'

Description: This stack creates a lambda for s3 trigger

Parameters:

BucketName:

Type: String

Default: 557bucket1

Resources:

Bucket:

Type: AWS::S3::Bucket

Properties:

BucketName: !Ref BucketName

NotificationConfiguration:

LambdaConfigurations:

- Event: 's3:ObjectCreated:*

Filter:

S3Key:

Rules:

- Name: suffix

Value: .csv

Function: !GetAtt Lambda.Arn

Lambda:

Type: AWS::Lambda::Function

Properties:

FunctionName: 'mynodelambda1'

Handler: index.handler

Role: arn:aws:iam::573351626142:role/service-role/mynodelambda-role-cd8ibrmh

Code:

S3Bucket: '558bucket1'

S3Key: 'code/lambdards.zip'

Runtime: "nodejs16.x"

Timeout: 300

MemorySize: 512

TracingConfig:

Mode: Active

S3InvokeLambdaPermission:

Type: AWS::Lambda::Permission

Properties:

Action: lambda:InvokeFunction

FunctionName: !Ref Lambda

Principal: s3.amazonaws.com

SourceArn: !Sub arn:aws:s3:::\${BucketName}

Lambda API to get called when a csv object is created in S3 Bucket persistent data in RDS (created in problem-1):

[Lambda API calls are batched in 100 to write 100 records from csv to the database if no of records of csv are greater than 1000]

```
const aws = require('aws-sdk');

const s3 = new aws.S3({
  apiVersion: '2006-03-01',
  region: 'us-east-1',
  accessKeyId: '#accesskey',
  secretAccessKey: '#secretKey'
});

const mysql = require('mysql');
const readline = require('readline');
let data = [];

const connection = mysql.createConnection({
  host: 'mydbinstance-us-east-1.cc7otgq3h7hn.us-east-1.rds.amazonaws.com',
  user: 'admin',
  password: 'admin1234',
  database: 'mydb',
  port: 3306
});

exports.handler = async (event, context, callback) => {
  // TODO implement
  var bucketName = event["Records"][0]["s3"]["bucket"]["name"];
  var bucketObject = event["Records"][0]["s3"]["object"]["key"];
  const bucketParams = {
    Bucket: `${bucketName}`,
    Key: `${bucketObject}`
```

```

};

console.log(bucketParams);

try {
    var tbname = bucketObject.split('.')[0];

    var crque = `create table ${tbname} (id int, name varchar(50),surname varchar(50), dob
varchar(50), gender varchar(10))`;

    console.log(crque);

    connection.query(crque, (err, res) => {
        if (err) {
            console.error(err);
        }
        console.log('table created');
    });
});

const s3ReadStream = s3.getObject(bucketParams).createReadStream();

const rl = readline.createInterface({
    input: s3ReadStream,
    terminal: false
});

var myReadPromise = new Promise((resolve, reject) => {
    var c = 0;

    rl.on('line', (line) => {
        c++;

        line = line.split(',');

        data.push(line);

        if (c % 100 == 0) {
            var i = (c-100);
            while (i < data.length) {

```

```

        var que = `insert into ${tbname}
values(${data[i][0]},'${data[i][1]}','${data[i][2]}','${data[i][3]}','${data[i][4]}')`;

        console.log(que);

        connection.query(que, (err, res) => {

            if (err) {

                console.error(err);

            }

        });

        i+=1;

    }

    console.log((i - 1) + ' records inserted');

}

});

rl.on('error', () => {

    console.log('error');

});

rl.on('close',()=>{

    var i=1;

    if(c<100){

        while (i < data.length) {

            var que = `insert into ${tbname}
values(${data[i][0]},'${data[i][1]}','${data[i][2]}','${data[i][3]}','${data[i][4]}')`;

            console.log(que);

            connection.query(que, (err, res) => {

                if (err) {

                    console.error(err);

                }

            });

            i+=1;

        }

    }

}

```

```

        console.log((i - 1) + ' records inserted');
    }
});

});

await myReadPromise;
} catch (err) {
    console.log(err);
}

};

```

Uploading a .csv file to the bucket created:

The screenshot displays the AWS Management Console interface for an Amazon S3 bucket named '557bucket1'. The 'Objects' tab is selected, showing a list of objects. A green notification banner at the top indicates 'Upload succeeded' with a link to 'View details below'.

Below the notification, the 'Upload: status' section provides a summary of the upload process:

Summary		
Destination s3://557bucket1	Succeeded 1 file, 123.0 B (100.00%)	Failed 0 files, 0 B (0%)

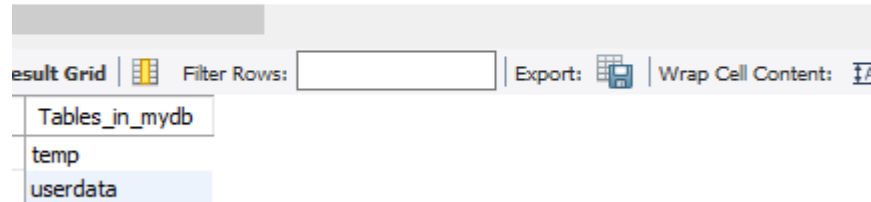
The 'Files and folders' section shows a total of 1 file and 123.0 B of data uploaded.

Verifying the database:

```
10 • show tables;
```

```
11
```

```
12 • select * from temp;
```

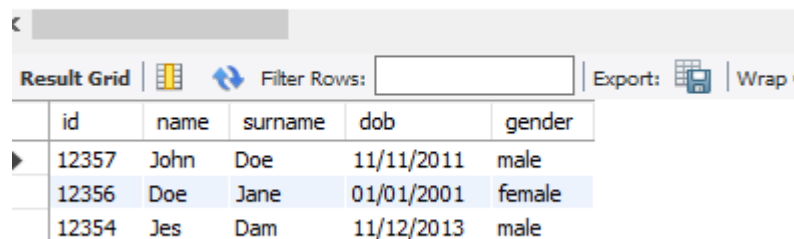


The screenshot shows a database interface with a table named 'Tables_in_mydb'. The table has two rows: 'temp' and 'userdata'. The 'temp' row is highlighted.

Tables_in_mydb
temp
userdata

```
11
```

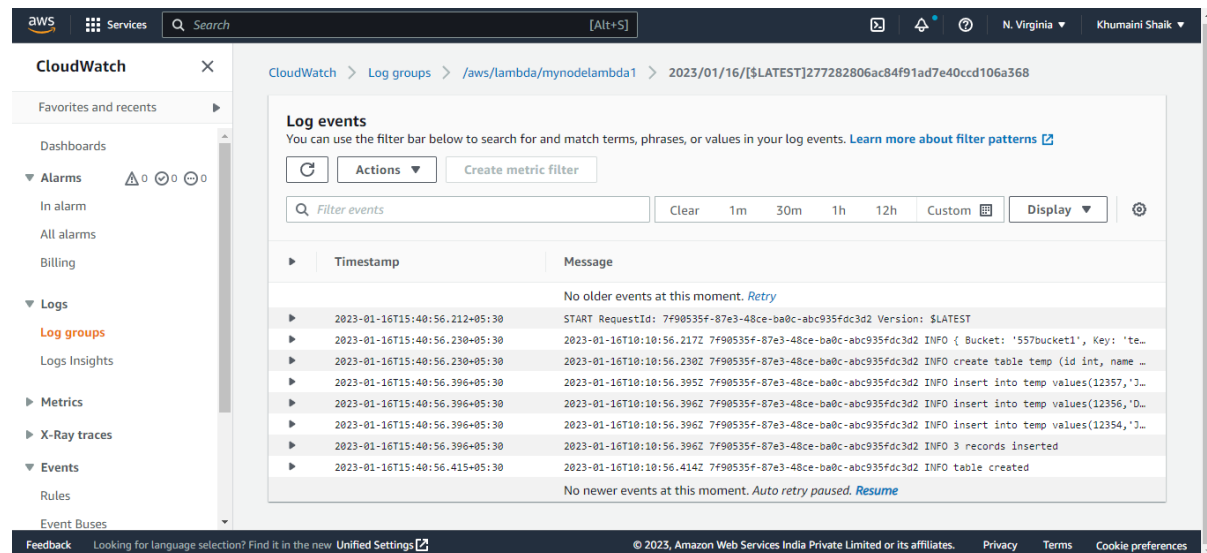
```
12 • select * from temp;
```



The screenshot shows a database interface with a table named 'temp'. The table has five columns: 'id', 'name', 'surname', 'dob', and 'gender'. There are three rows of data.

	id	name	surname	dob	gender
▶	12357	John	Doe	11/11/2011	male
	12356	Doe	Jane	01/01/2001	female
	12354	Jes	Dam	11/12/2013	male

Verifying the CloudWatch Logs to view the output of the Lambda function:



The screenshot shows the AWS CloudWatch console. The left sidebar contains navigation options: Dashboards, Alarms, Logs, Metrics, X-Ray traces, Events, and Event Buses. The main area displays the 'Log events' for the Lambda function '/aws/lambda/mynodelambda1'. The log events are filtered by the pattern '2023/01/16/[\$LATEST]277282806ac84f91ad7e40ccd106a368'. The log events show the following messages:

- 2023-01-16T15:40:56.212+05:30: No older events at this moment. [Retry](#)
- 2023-01-16T15:40:56.230+05:30: START RequestId: 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 Version: \$LATEST
- 2023-01-16T15:40:56.230+05:30: 2023-01-16T10:10:56.217Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO { Bucket: '557bucket1', Key: 'te...
- 2023-01-16T15:40:56.395Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO create table temp (id int, name ...
- 2023-01-16T15:40:56.396Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO insert into temp values(12357,'J...
- 2023-01-16T15:40:56.396Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO insert into temp values(12356,'D...
- 2023-01-16T15:40:56.396Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO insert into temp values(12354,'J...
- 2023-01-16T15:40:56.415+05:30: 2023-01-16T10:10:56.396Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO 3 records inserted
- 2023-01-16T15:40:56.415+05:30: 2023-01-16T10:10:56.414Z 7f90535f-87e3-48ce-ba0c-abc935fdc3d2 INFO table created
- 2023-01-16T15:40:56.415+05:30: No newer events at this moment. [Auto retry paused. Resume](#)